

```

# Install necessary dependencies
!pip install -q opencv-python-headless
!pip install -q matplotlib
!pip install -q numpy
!pip install -q scikit-image
!pip install -q tensorflow
!pip install -q albumentations

import os
import cv2
import numpy as np
import matplotlib.pyplot as plt
from google.colab import files
from sklearn.cluster import KMeans
from skimage import exposure, filters, color
import random
from typing import List, Tuple, Dict, Any, Optional
import glob
import time

# Set random seeds for reproducibility
np.random.seed(42)
random.seed(42)

class CarDamagePreprocessor:
    """
    A class for preprocessing damaged car images to prepare them for damage detection
    and classification models.
    """

    def __init__(self,
                 target_size: Tuple[int, int] = (512, 512),
                 normalize: bool = True,
                 clahe_clip_limit: float = 2.0,
                 clahe_grid_size: Tuple[int, int] = (8, 8)):

        self.target_size = target_size
        self.normalize = normalize
        self.clahe = cv2.createCLAHE(clipLimit=clahe_clip_limit,
                                     tileGridSize=clahe_grid_size)

    def load_image(self, image_path: str) -> np.ndarray:
        """
        The loaded image as a numpy array
        """
        dardize_image(self, image: np.ndarray) -> np.ndarray:

        # Resize to target size
        resized = cv2.resize(image, (self.target_size[1], self.target_size[0]))

        # Normalize pixel values if requested
        if self.normalize:
            return resized.astype(np.float32) / 255.0
        return resized

    def remove_background(self, image: np.ndarray,
                        threshold: int = 25,
                        blur_size: int = 5) -> Tuple[np.ndarray, np.ndarray]:

        # Create a copy of the image
        img = image.copy()

        # Convert to RGB if needed
        if len(img.shape) == 2:
            img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)

        # Ensure image is uint8 for GrabCut (required by OpenCV)
        if img.dtype == np.float32:
            img = (img * 255).astype(np.uint8)
        elif img.dtype != np.uint8:
            img = img.astype(np.uint8)

        # Initial mask creation
        mask = np.zeros(img.shape[:2], np.uint8)

        # Background and foreground models
        bgd_model = np.zeros((1, 65), np.float64)
        fgd_model = np.zeros((1, 65), np.float64)

        # Define rough ROI around the image center assuming car is in the middle
        margin = 50

```

```

rect = (margin, margin, img.shape[1]-2*margin, img.shape[0]-2*margin)

try:
    # Apply GrabCut
    cv2.grabCut(img, mask, rect, bgd_model, fgd_model, 5, cv2.GC_INIT_WITH_RECT)

    # Convert mask
    mask2 = np.where((mask==2) | (mask==0), 0, 1).astype('uint8')
except cv2.error:
    # Fallback if GrabCut fails
    print("GrabCut failed. Using basic thresholding as fallback.")
    gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
    blurred = cv2.GaussianBlur(gray, (blur_size, blur_size), 0)
    _, mask2 = cv2.threshold(blurred, threshold, 1, cv2.THRESH_BINARY)

# Apply the mask to the image
result = img * mask2[:, :, np.newaxis]

# Convert back to original format if needed
if image.dtype == np.float32:
    result = result.astype(np.float32) / 255.0

return result, mask2

def detect_roi(self, image: np.ndarray,
               mask: Optional[np.ndarray] = None) -> Tuple[np.ndarray, Tuple[int, int, int, int]]:

    # If mask is provided, use it to find contours
    if mask is not None:
        contours, _ = cv2.findContours(mask, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
        if contours:
            # Find the largest contour (assumed to be the car)
            largest_contour = max(contours, key=cv2.contourArea)
            x, y, w, h = cv2.boundingRect(largest_contour)

            # Add some padding
            padding = 10
            x = max(0, x - padding)
            y = max(0, y - padding)
            w = min(image.shape[1] - x, w + 2*padding)
            h = min(image.shape[0] - y, h + 2*padding)

            # Crop the image to the bounding box
            cropped = image[y:y+h, x:x+w]
            return cropped, (x, y, w, h)

    # If no mask or no contours found, use edge detection as fallback
    gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY) if len(image.shape) > 2 else image
    blurred = cv2.GaussianBlur(gray, (5, 5), 0)
    edges = cv2.Canny(blurred, 50, 150)

    # Find contours in the edge map
    contours, _ = cv2.findContours(edges, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)

    if contours:
        # Combine all contours to find the overall bounding box
        all_points = np.concatenate([cnt for cnt in contours])
        x, y, w, h = cv2.boundingRect(all_points)

        # Add some padding
        padding = 20
        x = max(0, x - padding)
        y = max(0, y - padding)
        w = min(image.shape[1] - x, w + 2*padding)
        h = min(image.shape[0] - y, h + 2*padding)

        # Crop the image to the bounding box
        cropped = image[y:y+h, x:x+w]
        return cropped, (x, y, w, h)

    # If all else fails, return the original image
    return image, (0, 0, image.shape[1], image.shape[0])

def reduce_noise(self, image: np.ndarray,
                 method: str = 'gaussian',
                 kernel_size: int = 5) -> np.ndarray:

    if method == 'gaussian':
        return cv2.GaussianBlur(image, (kernel_size, kernel_size), 0)
    elif method == 'median':
        return cv2.medianBlur(image, kernel_size)
    elif method == 'bilateral':
        if len(image.shape) > 2 and image.dtype == np.float32:

```

```

        # Convert to 8-bit for bilateral filter
        temp = (image * 255).astype(np.uint8)
        result = cv2.bilateralFilter(temp, kernel_size, 75, 75)
        return result.astype(np.float32) / 255.0
    else:
        return cv2.bilateralFilter(image, kernel_size, 75, 75)
else:
    raise ValueError(f"Unknown noise reduction method: {method}")

def enhance_contrast(self, image: np.ndarray,
                    method: str = 'clahe') -> np.ndarray:

    # Convert to grayscale if image is RGB
    if len(image.shape) > 2:
        gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
    else:
        gray = image.copy()

    # Scale to 0-255 if normalized
    if gray.dtype == np.float32:
        gray = (gray * 255).astype(np.uint8)

    if method == 'clahe':
        enhanced = self.clahe.apply(gray)
    elif method == 'histeq':
        enhanced = cv2.equalizeHist(gray)
    elif method == 'adapthist':
        enhanced = exposure.equalize_adapthist(gray, clip_limit=0.03)
        enhanced = (enhanced * 255).astype(np.uint8)
    else:
        raise ValueError(f"Unknown contrast enhancement method: {method}")

    # If input was RGB, convert back to RGB
    if len(image.shape) > 2:
        # Create a 3-channel image where each channel has the enhanced data
        enhanced_rgb = np.zeros_like(image)
        if image.dtype == np.float32:
            enhanced_rgb[:, :, 0] = enhanced.astype(np.float32) / 255.0
            enhanced_rgb[:, :, 1] = enhanced.astype(np.float32) / 255.0
            enhanced_rgb[:, :, 2] = enhanced.astype(np.float32) / 255.0
        else:
            enhanced_rgb[:, :, 0] = enhanced
            enhanced_rgb[:, :, 1] = enhanced
            enhanced_rgb[:, :, 2] = enhanced
        return enhanced_rgb

    # Return the enhanced grayscale image
    if image.dtype == np.float32:
        return enhanced.astype(np.float32) / 255.0
    return enhanced

def detect_edges(self, image: np.ndarray,
                method: str = 'canny',
                low_threshold: int = 50,
                high_threshold: int = 150) -> np.ndarray:

    # Convert to grayscale if image is RGB
    if len(image.shape) > 2:
        gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
    else:
        gray = image.copy()

    # Scale to 0-255 if normalized
    if gray.dtype == np.float32:
        gray = (gray * 255).astype(np.uint8)

    # Apply Gaussian blur to reduce noise
    blurred = cv2.GaussianBlur(gray, (5, 5), 0)

    if method == 'canny':
        edges = cv2.Canny(blurred, low_threshold, high_threshold)
    elif method == 'sobel':
        sobelx = cv2.Sobel(blurred, cv2.CV_64F, 1, 0, ksize=3)
        sobely = cv2.Sobel(blurred, cv2.CV_64F, 0, 1, ksize=3)
        edges = np.sqrt(sobelx**2 + sobely**2)
        edges = cv2.normalize(edges, None, 0, 255, cv2.NORM_MINMAX, cv2.CV_8U)
    elif method == 'scharr':
        scharrx = cv2.Scharr(blurred, cv2.CV_64F, 1, 0)
        scharry = cv2.Scharr(blurred, cv2.CV_64F, 0, 1)
        edges = np.sqrt(scharrx**2 + scharry**2)
        edges = cv2.normalize(edges, None, 0, 255, cv2.NORM_MINMAX, cv2.CV_8U)
    else:
        raise ValueError(f"Unknown edge detection method: {method}")

```

```

    raise ValueError('Unknown edge detection method: {method}'.format(method=method))

    # Return the edge map
    if image.dtype == np.float32:
        return edges.astype(np.float32) / 255.0
    return edges

def segment_damage(self, image: np.ndarray,
                   edge_map: np.ndarray = None,
                   threshold: float = 0.3) -> np.ndarray:

    if edge_map is None:
        edge_map = self.detect_edges(image)

    # Threshold the edge map to get binary mask
    if edge_map.dtype == np.float32:
        mask = (edge_map > threshold).astype(np.uint8)
    else:
        mask = (edge_map > threshold * 255).astype(np.uint8)

    # Apply morphological operations to clean up the mask
    kernel = np.ones((5, 5), np.uint8)
    mask = cv2.morphologyEx(mask, cv2.MORPH_CLOSE, kernel)
    mask = cv2.morphologyEx(mask, cv2.MORPH_OPEN, kernel)

    # Label connected components
    num_labels, labels = cv2.connectedComponents(mask)

    # Filter out small regions
    min_size = 50
    for i in range(1, num_labels):
        if np.sum(labels == i) < min_size:
            mask[labels == i] = 0

    return mask

def extract_features(self, image: np.ndarray,
                    mask: Optional[np.ndarray] = None) -> Dict[str, Any]:

    # Convert to grayscale if image is RGB
    if len(image.shape) > 2:
        gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
    else:
        gray = image.copy()

    # Apply mask if provided
    if mask is not None:
        masked_gray = cv2.bitwise_and(gray, gray, mask=mask)
    else:
        masked_gray = gray

    # Scale to 0-255 if normalized
    if masked_gray.dtype == np.float32:
        masked_gray = (masked_gray * 255).astype(np.uint8)

    # Extract features
    features = {}

    # Basic statistics
    if np.any(masked_gray > 0):
        features['mean'] = np.mean(masked_gray[masked_gray > 0])
        features['std'] = np.std(masked_gray[masked_gray > 0])
        features['min'] = np.min(masked_gray[masked_gray > 0])
        features['max'] = np.max(masked_gray[masked_gray > 0])
    else:
        features['mean'] = 0
        features['std'] = 0
        features['min'] = 0
        features['max'] = 0

    # Histogram
    hist = cv2.calcHist([masked_gray], [0], None, [256], [0, 256])
    features['histogram'] = hist.flatten()

    # Texture features using Haralick texture features (calculated manually)
    if np.any(masked_gray > 0):
        # Convert to uint8 for texture analysis
        masked_gray_uint8 = masked_gray.astype(np.uint8)

        # Calculate gradient magnitude as a simple texture feature
        sobelx = cv2.Sobel(masked_gray_uint8, cv2.CV_64F, 1, 0, ksize=3)
        sobely = cv2.Sobel(masked_gray_uint8, cv2.CV_64F, 0, 1, ksize=3)
        gradient_magnitude = np.sqrt(sobelx**2 + sobely**2)

```

```

features['gradient_mean'] = np.mean(gradient_magnitude)
features['gradient_std'] = np.std(gradient_magnitude)

# Calculate local binary pattern (simple version)
def local_binary_pattern(image, points=8, radius=1):
    rows, cols = image.shape
    result = np.zeros((rows-2*radius, cols-2*radius), dtype=np.uint8)
    for i in range(radius, rows-radius):
        for j in range(radius, cols-radius):
            center = image[i, j]
            pattern = 0
            for p in range(points):
                angle = 2 * np.pi * p / points
                x = j + int(round(radius * np.cos(angle)))
                y = i + int(round(radius * np.sin(angle)))
                if image[y, x] >= center:
                    pattern |= (1 << p)
            result[i-radius, j-radius] = pattern
    return result

try:
    # Only compute LBP on a smaller region if image is large
    if masked_gray_uint8.shape[0] > 100 and masked_gray_uint8.shape[1] > 100:
        center_y, center_x = masked_gray_uint8.shape[0] // 2, masked_gray_uint8.shape[1] // 2
        roi_size = 50
        roi = masked_gray_uint8[
            max(0, center_y - roi_size):min(masked_gray_uint8.shape[0], center_y + roi_size),
            max(0, center_x - roi_size):min(masked_gray_uint8.shape[1], center_x + roi_size)
        ]
        lbp = local_binary_pattern(roi)
    else:
        lbp = local_binary_pattern(masked_gray_uint8)

    lbp_hist = cv2.calcHist([lbp], [0], None, [256], [0, 256])
    features['lbp_histogram'] = lbp_hist.flatten()
    features['lbp_entropy'] = -np.sum((lbp_hist / np.sum(lbp_hist)) *
                                     np.log2(lbp_hist / np.sum(lbp_hist) + 1e-10))

except Exception as e:
    print(f"LBP calculation error: {e}")
    features['lbp_histogram'] = np.zeros(256)
    features['lbp_entropy'] = 0

else:
    features['gradient_mean'] = 0
    features['gradient_std'] = 0
    features['lbp_histogram'] = np.zeros(256)
    features['lbp_entropy'] = 0

# SIFT features (keypoints)
try:
    if np.any(masked_gray > 0):
        sift = cv2.SIFT_create()
        keypoints, descriptors = sift.detectAndCompute(masked_gray, None)
        features['num_keypoints'] = len(keypoints)
        features['keypoints'] = keypoints
        features['descriptors'] = descriptors if descriptors is not None else np.array([])
    else:
        features['num_keypoints'] = 0
        features['keypoints'] = []
        features['descriptors'] = np.array([])
except Exception as e:
    print(f"SIFT feature extraction error: {e}")
    features['num_keypoints'] = 0
    features['keypoints'] = []
    features['descriptors'] = np.array([])

return features

def augment_data(self, image: np.ndarray,
                 num_augmentations: int = 5) -> List[np.ndarray]:

    augmented_images = []

    # Define some augmentation functions
    def random_brightness_contrast(img, brightness_range=(-0.2, 0.2), contrast_range=(-0.2, 0.2)):
        # Brightness adjustment
        brightness = np.random.uniform(brightness_range[0], brightness_range[1])
        adjusted = img.astype(np.float32) + brightness

        # Contrast adjustment
        contrast = np.random.uniform(contrast_range[0], contrast_range[1]) + 1.0
        adjusted = adjusted * contrast

```

```

    # Clip values to valid range
    adjusted = np.clip(adjusted, 0, 1.0 if img.dtype == np.float32 else 255)
    return adjusted.astype(img.dtype)

def random_noise(img, var=0.01):
    # Add Gaussian noise
    if img.dtype == np.float32:
        noise = np.random.normal(0, var**0.5, img.shape)
        noisy = img + noise
        return np.clip(noisy, 0, 1.0).astype(np.float32)
    else:
        noise = np.random.normal(0, var**0.5 * 255, img.shape).astype(np.int16)
        noisy = img.astype(np.int16) + noise
        return np.clip(noisy, 0, 255).astype(np.uint8)

def random_rotation(img, angle_range=(-15, 15)):
    # Random rotation
    angle = np.random.uniform(angle_range[0], angle_range[1])
    rows, cols = img.shape[:2]
    M = cv2.getRotationMatrix2D((cols/2, rows/2), angle, 1)
    return cv2.warpAffine(img, M, (cols, rows))

def random_flip(img):
    # Random horizontal flip
    if np.random.random() > 0.5:
        return cv2.flip(img, 1)
    return img

def random_crop(img, crop_factor_range=(0.8, 0.95)):
    factor = np.random.uniform(crop_factor_range[0], crop_factor_range[1])
    h, w = img.shape[:2]
    crop_h, crop_w = int(h * factor), int(w * factor)
    start_h = np.random.randint(0, h - crop_h + 1)
    start_w = np.random.randint(0, w - crop_w + 1)
    cropped = img[start_h:start_h+crop_h, start_w:start_w+crop_w]
    return cv2.resize(cropped, (w, h))

# Define augmentation pipeline with probabilities
augmentation_functions = [
    (random_brightness_contrast, 0.7),
    (random_noise, 0.5),
    (random_rotation, 0.5),
    (random_flip, 0.5),
    (random_crop, 0.5)
]

for _ in range(num_augmentations):
    # Start with a copy of the original image
    augmented = image.copy()

    # Apply random augmentations based on probability
    for aug_func, prob in augmentation_functions:
        if np.random.random() < prob:
            augmented = aug_func(augmented)

    augmented_images.append(augmented)

return augmented_images

def visualize_preprocessing(self, original: np.ndarray,
                           processed_results: Dict[str, np.ndarray]) -> None:

    # Determine number of steps
    n_steps = len(processed_results) + 1 # +1 for original

    # Create figure with subplots
    fig, axes = plt.subplots(1, n_steps, figsize=(20, 5))

    # Plot original image
    axes[0].imshow(original)
    axes[0].set_title('Original')
    axes[0].axis('off')

    # Plot processed results
    for i, (title, img) in enumerate(processed_results.items(), 1):
        # Handle different image types
        if len(img.shape) == 2: # Grayscale or mask
            if img.dtype == bool:
                img = img.astype(np.uint8) * 255

            # Display as grayscale
            axes[i].imshow(img, cmap='gray')
        else:

```

```

-----
    # Display as RGB
    if img.dtype == np.float32 and np.max(img) <= 1.0:
        axes[i].imshow(img)
    else:
        axes[i].imshow(img.astype(np.uint8))

    axes[i].set_title(title)
    axes[i].axis('off')

plt.tight_layout()
plt.show()

def process_image(self, image: np.ndarray,
                  visualize: bool = False) -> Dict[str, Any]:

    results = {}

    # Standardize image
    std_image = self.standardize_image(image)
    results['standardized'] = std_image

    # Remove background
    bg_removed, mask = self.remove_background(std_image)
    results['background_removed'] = bg_removed
    results['background_mask'] = mask

    # Detect ROI
    roi, bbox = self.detect_roi(bg_removed, mask)
    results['roi'] = roi
    results['bbox'] = bbox

    # Reduce noise
    denoised = self.reduce_noise(roi, method='bilateral')
    results['denoised'] = denoised

    # Enhance contrast
    enhanced = self.enhance_contrast(denoised)
    results['enhanced'] = enhanced

    # Detect edges
    edges = self.detect_edges(enhanced)
    results['edges'] = edges

    # Segment damage
    damage_mask = self.segment_damage(enhanced, edges)
    results['damage_mask'] = damage_mask

    # Extract features
    features = self.extract_features(enhanced, damage_mask)
    results['features'] = features

    # Visualize if requested
    if visualize:
        vis_results = {
            'Background Removed': bg_removed,
            'ROI': roi,
            'Denoised': denoised,
            'Enhanced': enhanced,
            'Edges': edges,
            'Damage Mask': damage_mask
        }
        self.visualize_preprocessing(image, vis_results)

    return results

def process_directory(self, directory_path: str,
                     output_dir: str = None,
                     visualize: bool = False) -> Dict[str, Dict[str, Any]]:

    # Create output directory if specified
    if output_dir is not None:
        os.makedirs(output_dir, exist_ok=True)

    results = {}

    # Get all image files
    image_files = []
    for ext in ['*.jpg', '*.jpeg', '*.png', '*.bmp']:
        image_files.extend(glob.glob(os.path.join(directory_path, ext)))
        image_files.extend(glob.glob(os.path.join(directory_path, ext.upper())))

    print(f"Found {len(image_files)} images in {directory_path}")

```

```

# Process each image
for image_file in image_files:
    try:
        # Load image
        image = self.load_image(image_file)

        # Process image
        result = self.process_image(image, visualize=visualize)

        # Save processed images if output directory is specified
        if output_dir is not None:
            # Get base filename without extension
            basename = os.path.splitext(os.path.basename(image_file))[0]

            # Save each processed image
            for name, img in result.items():
                if isinstance(img, np.ndarray):
                    # Create image file path
                    img_path = os.path.join(output_dir, f"{basename}_{name}.png")

                    # Convert to uint8 if needed
                    if img.dtype == np.float32:
                        img = (img * 255).astype(np.uint8)

                    # Save the image
                    if len(img.shape) == 2:
                        cv2.imwrite(img_path, img)
                    else:
                        cv2.imwrite(img_path, cv2.cvtColor(img, cv2.COLOR_RGB2BGR))

            # Store results
            results[os.path.basename(image_file)] = result

    except Exception as e:
        print(f"Error processing {image_file}: {e}")

return results

# Example usage
def main():
    """
    Example usage of the CarDamagePreprocessor.
    """
    # Create preprocessor
    preprocessor = CarDamagePreprocessor()

    # Check if images already exist in the environment
    import os
    existing_images = [f for f in os.listdir() if f.lower().endswith(('.png', '.jpg', '.jpeg', '.bmp'))]

    if existing_images:
        print(f"Found {len(existing_images)} images in the current directory.")
        image_files = existing_images
    else:
        # Allow user to upload images
        print("Please upload one or more damaged car images.")
        uploaded = files.upload()
        image_files = list(uploaded.keys())

    # Process images
    for filename in image_files:
        try:
            print(f"Processing {filename}...")

            # Load image
            image = cv2.imread(filename)
            if image is None:
                print(f"Error: Could not read image {filename}")
                continue

            image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

            # Process image with error handling
            try:
                result = preprocessor.process_image(image, visualize=True)

                # Display features
                print(f"Image features:")
                for key, value in result['features'].items():
                    if key in ['histogram', 'keypoints', 'descriptors']:
                        if isinstance(value, np.ndarray):
                            print(f"  {key}: [array with shape {value.shape}]")

```



```
        else:
            print(f" {key}: [array with {len(value)} elements]")
    else:
        print(f" {key}: {value}")

print("\n")

# Demonstrate augmentation with the first successful image
print("Generating data augmentations...")
augmented_images = preprocessor.augment_data(image, num_augmentations=5)

# Display augmented images
plt.figure(figsize=(15, 10))
plt.subplot(2, 3, 1)
plt.imshow(image)
plt.title("Original")
plt.axis('off')

for i, aug_img in enumerate(augmented_images, 1):
    plt.subplot(2, 3, i+1)
    plt.imshow(aug_img)
    plt.title(f"Augmentation {i}")
    plt.axis('off')

plt.tight_layout()
plt.show()

# Only process one image for demonstration
break

except Exception as e:
    print(f"Error during image processing: {str(e)}")
    import traceback
    traceback.print_exc()

except Exception as e:
    print(f"Error with image {filename}: {str(e)}")
    import traceback
    traceback.print_exc()

if __name__ == "__main__":
    main()
```

